

# Seminar notebook

Anastasia Chanbour

```
knitr::opts_chunk$set(echo = TRUE)
library(tidyverse)
library(broom) #broom package takes the messy output of built-in functions in R, such as lm, #nls, or t.
theme_set(theme_minimal(base_size = 10)) #theme_set overrides default ggplot theme
library(gapminder)
library(plotly)
```

## Linear regression

### Simple linear regression

#### Estimation

**Task** 1. Filter year=2007 from the gapminder data and call the data frame object gm2007. 2. Create a regression object called model\_lm that regresses lifeExp on gdpPercap for 2007 data 3. Output the estimated coefficients, intercept and slope, only

```
#CODE
gm2007 <- gapminder %>% filter(year == 2007)
model_lm <- lm(data = gm2007, lifeExp ~ gdpPercap)
#tidy(model_lm)['estimate']
coef(model_lm)
```

```
## (Intercept)    gdpPercap
## 5.956565e+01 6.371341e-04
```

#### Demo dplyr::summarise function

summarise() or summarize() creates a new data frame.

dplyr = (d)ata + plyr plyr: Tools for Splitting, Applying and Combining Data

- $\hat{\beta}_1 = \text{cor}(x, y) * \text{sd}(y) / \text{sd}(x)$
- $\hat{\beta}_0 = \text{mean}(y) - \hat{\beta}_1 * \text{mean}(x)$
- Regression line passes through  $(\text{mean}(x), \text{mean}(y))$

*#output: correlation, sd, mean(x), mean(y), estimated intercept and slope coefficient using formula*  
*#CODE*

```
gm2007 %>%
  summarize(cor_xy=cor(gdpPercap, lifeExp),
            sd_x = sd(gdpPercap),
            sd_y = sd(lifeExp),
            mean_x = mean(gdpPercap),
            mean_y = mean(lifeExp),
            hat_beta1 = cor_xy/sd_x*sd_y,
            hat_beta0 = mean_y - hat_beta1*mean_x)
```

```
## # A tibble: 1 x 7
##   cor_xy   sd_x   sd_y mean_x mean_y hat_beta1 hat_beta0
##   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl>
## 1  0.679 12860.  12.1 11680.   67.0 0.000637   59.6
```

## Inference

*#Use summary(reg\_object) to obtain basic output of regression*  
summary(model\_lm)

```
##
## Call:
## lm(formula = lifeExp ~ gdpPercap, data = gm2007)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -22.828  -6.316   1.922   6.898  13.128
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 5.957e+01  1.010e+00  58.95  <2e-16 ***
## gdpPercap   6.371e-04  5.827e-05  10.93  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.899 on 140 degrees of freedom
## Multiple R-squared:  0.4606, Adjusted R-squared:  0.4567
## F-statistic: 119.5 on 1 and 140 DF, p-value: < 2.2e-16
```

(ISLR eq 3.8) p.66

```
se(beta hat) = sigma / sqrt(sum((x - mean(x))^2))
```

Estimated by:

```
se(beta hat) = sigma hat / sqrt(sum((x - mean(x))^2))
```

where (ISLR 3.15) p.69:

```
sigma hat = RSE = sqrt( RSS / (n-2) )
```

*#Using augment() to extract residuals, compute the estimated se of the slope parameter by  
#obtaining 1. RSS 2. MSE = (RSS/(n-p)) 3. estimated S.E of hat beta1*

*#CODE*

```
augment(model_lm) %>%
  summarize(RSS = sum(.resid^2),
            MSE = RSS / (n()-2),
            SE = sqrt(MSE)/sqrt(sum((gdpPercap - mean(gdpPercap))^2)))
```

```
## # A tibble: 1 x 3
##   RSS    MSE      SE
##   <dbl> <dbl>   <dbl>
## 1 11086.  79.2 0.0000583
```

## Model diagnostics

```
#glance: Construct a single row summary "glance" of a model, fit, or other object
#CODE
glance(model_lm)
```

```
## # A tibble: 1 x 12
##   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
##   <dbl>      <dbl> <dbl>      <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    0.461        0.457  8.90        120. 1.69e-20     1 -511. 1028. 1037.
## # ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

“Portion of variance in outcome **explained** by simple linear regression model”

$$R^2 = \text{cor}(x, y)^2$$

```
#compute correlation squared between gdpPercap and lifeExp
#CODE
cor(gm2007$gdpPercap, gm2007$lifeExp)^2
```

```
## [1] 0.4605827
```

$$R^2 = 1 - \frac{\text{RSS}}{\text{TSS}}$$

```
#using residuals from augment(), obtain values of RSS, TSS, and R2
#CODE
augment(model_lm) %>%
  summarize(RSS = sum(.resid^2),
            TSS = sum((lifeExp - mean(lifeExp))^2),
            R2 = 1 - RSS/TSS)
```

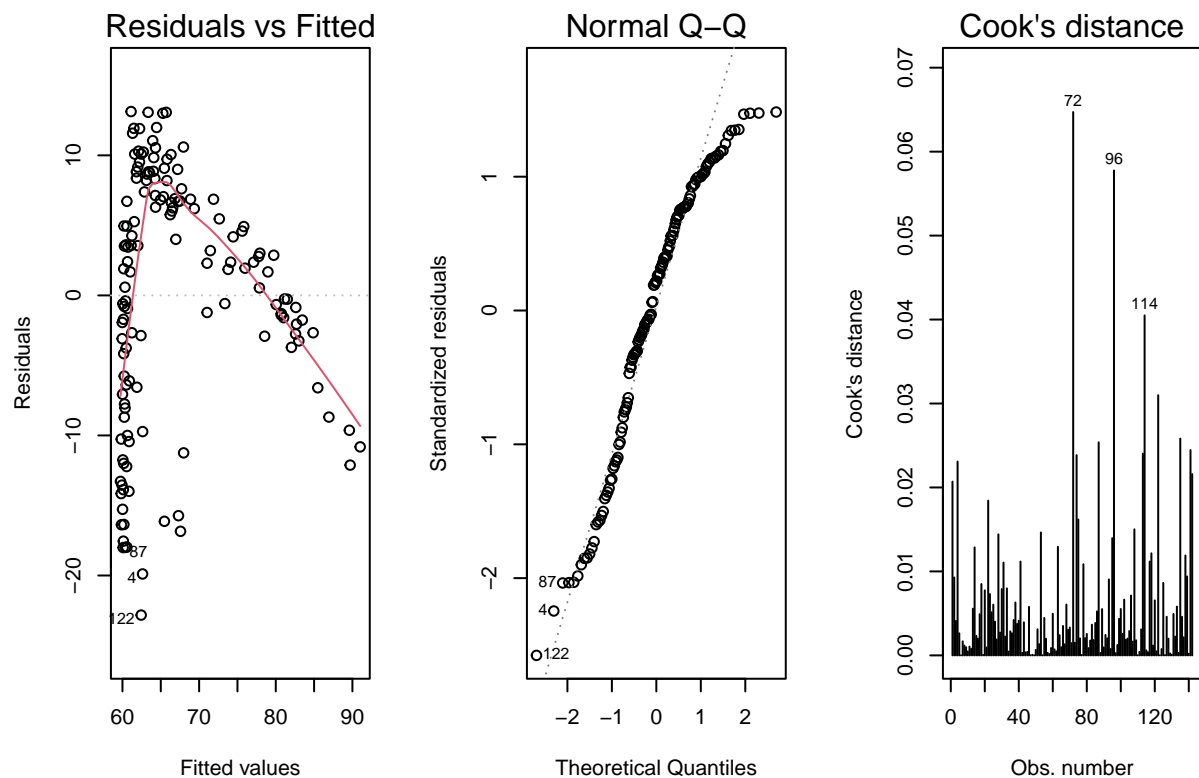
```
## # A tibble: 1 x 3
##   RSS    TSS    R2
##   <dbl> <dbl> <dbl>
## 1 11086. 20552. 0.461
```

## Diagnostic plots

Idea: look for patterns in residuals, which could indicate systematic errors (bias)

Outliers vs influential points: influential points affect the slope of the model line

```
#Plot residuals vs gdpPercap, using residuals from augment()
#CODE
par(mfrow=c(1,3))
plot(model_lm, which = c(1,2,4)) #plot.lm is plot on a regression generates 6 plots
```



There is a pattern in residuals. Suggests that the relationship is non-linear

Other diagnostics:

- Checking for (approximate) normality with quantile-quantile plot
- Checking for influential observations

Cook's distance, `cooks` in the plots, measures how much the predictions for all other observations change if we leave out one observation

Point with high `cooks` values