

Generalisation

Anastasia Chanbour

29/10/2021

Loading the required packages

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr 0.3.4
## v tibble 3.1.5       v dplyr 1.0.7
## v tidyr 1.1.4        v stringr 1.4.0
## v readr 2.0.2        v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(broom)
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
##
## Loaded glmnet 4.1-3
```

High-dimensional regression

Generating some data from a high-dimensional model (increasing p below)

```
set.seed(1)
n <- 100
p <- 3
X <- matrix(rnorm(n*p), nrow = n)
beta <- rpois(p, lambda = 1)
y <- X %*% beta + rnorm(n)
```

Ridge regression with glmnet


```
model_ridge <- glmnet(X, y,
                      intercept = FALSE,
```

```

alpha = 0,
lambda = 10^(-3:3))

plot(model_ridge, xvar = "lambda")

```



Estimation error

Computing the mean-squared error of the coefficient estimates at different values of lambda

```

lambda <- 10^0
beta_hat <- coef(model_ridge, s = lambda)[-1] # leave out intercept
mean((beta - beta_hat)^2)

```

```
## [1] 0.3569038
```

Prediction error

Computing predictions using the estimated coefficients and the mean-squared prediction error at different values of lambda

```

# e.g. predict(model_ridge, newx = X, s = lambda or 10*lambda)
y_hat <- X%*%beta_hat

#mse prediction of yhat
mean((y-y_hat)^2)

```

```
## [1] 1.3833
```

Overfitting to variance and In Distribution (ID) generalisation

Generating a new sample from the same distribution (which things stay fixed?)

```
X_ID <- matrix(rnorm(n*p),nrow=n)
y_ID <- X_ID%%beta + rnorm(n)
```

Calculating the prediction error on this new sample at different values of lambda

```
y_ID_hat <- predict(model_ridge, newx = X_ID, s = lambda )
#MSE
mean((y_ID - y_ID_hat)^2)
```

```
## [1] 1.879422
```

“Overfitting to bias” and Out Of Distribution (OOD) generalisation

There are many ways to change the distribution for a new sample

- Changing beta (adding a small amount of noise)
- Changing the CEF some other way (e.g. adding a non-linear term)
- Changing the distribution of X and/or the errors

We choose adding a noise to the beta coefficient

```
beta_new_D <- beta + rnorm(p, sd=0.1)
X_OOD <- matrix(rnorm(n*p),nrow=n)
y_OOD <- X_OOD %%% beta_new_D + rnorm(n)

y_OOD_hat <- predict(model_ridge, newx = X_OOD,s=lambda)

#MSE
mean((y_OOD - y_OOD_hat)^2)
```

```
## [1] 1.942832
```

Comparison with gradient descent

```
least_squares_gradient <- function(x, y, beta) {
  -2 * t(x) %%% (y - x %%% beta) #+ 2 * beta
}

least_squares_loss <- function(x, y, beta) {
  sum((y - x %%% beta)^2)
}

beta_prev2 <- rep(0, p) # or random start point
grad_prev2 <- least_squares_gradient(X, y, beta_prev2)
beta_prev1 <- beta_prev2 + 0.1 * grad_prev2 / sqrt(sum(grad_prev2^2))
grad_prev1 <- least_squares_gradient(X, y, beta_prev1)
previous_loss <- least_squares_loss(X, y, beta_prev2)
next_loss <- least_squares_loss(X, y, beta_prev1)
steps <- 1
while (abs(previous_loss - next_loss) > 0.001) {
  grad_diff <- grad_prev1 - grad_prev2
  step_BB <- sum((beta_prev1 - beta_prev2) * grad_diff) / sum(grad_diff^2)
  # Barzilai-Borwein step size
```

```

beta_prev2 <- beta_prev1
beta_prev1 <- beta_prev1 - step_BB * grad_prev1

grad_prev2 <- grad_prev1
grad_prev1 <- least_squares_gradient(X, y, beta_prev1)

previous_loss <- next_loss
next_loss <- least_squares_loss(X, y, beta_prev1)

print(previous_loss)
steps <- steps + 1
}

```

```

## [1] 1128.978
## [1] 88.81974
## [1] 84.38873
## [1] 84.2721
## [1] 84.27054

```

```

beta_final <- beta_prev1

```